

Classification, Halfspaces, the Perceptron algorithm

Chris Cornwell

Feb 25, 2025

Outline

Classification tasks

Half-space model

Perceptron algorithm

Outline

Classification tasks

Half-space model

Perceptron algorithm

Example of Classification

Use some model to determine a digit that was (hand)written in an image

0, 1, 2, 3, 4, 5, 6, 7, 8, or 9.

Example of Classification

Use some model to determine a digit that was (hand)written in an image

0, 1, 2, 3, 4, 5, 6, 7, 8, or 9.

- ▶ Convert image to a vector (*in some way*) $\rightarrow \mathbf{x}$.
- ▶ Your model's output: $\hat{y}(\mathbf{x})$ is the (predicted) digit.

Provided with your data, an “observation” $y \in \{0, 1, \dots, 9\}$ of the digit being written.

Example of Classification


Use some model to determine a digit that was (hand)written in an image

0, 1, 2, 3, 4, 5, 6, 7, 8, or 9.

- ▶ Convert image to a vector (*in some way*) $\rightarrow \mathbf{x}$.
- ▶ Your model's output: $\hat{y}(\mathbf{x})$ is the (predicted) digit.

Provided with your data, an “observation” $y \in \{0, 1, \dots, 9\}$ of the digit being written.

y and \hat{y} are numbers on number line; but, use them like labels (or, separate buckets) to group points \mathbf{x} . When $y = 5$, predicting $\hat{y} = 4$ is not any better than $\hat{y} = 0$.

A handwritten digit '5' followed by a handwritten digit '0'.

Close only counts in horseshoes ...Regression

In linear regression, on indpt. variables x_0, x_1, \dots, x_{d-1} , had (affine) linear function $\hat{y} = p_0x_0 + p_1x_1 + \dots + p_{d-1}x_{d-1} + p_d$;
values of function \leftrightarrow prediction \hat{y} ; error term ε , so that $y = \hat{y} + \varepsilon$.

¹Should consider the output y here to be a random variable, with distribution that depends on x . In simple linear regression, ε is a random variable and

$$y = p_0x_0 + p_1 + \varepsilon.$$

Close only counts in horseshoes ...Regression

In linear regression, on indpt. variables x_0, x_1, \dots, x_{d-1} , had (affine) linear function $\hat{y} = p_0x_0 + p_1x_1 + \dots + p_{d-1}x_{d-1} + p_d$; values of function \leftrightarrow prediction \hat{y} ; error term ε , so that $y = \hat{y} + \varepsilon$. An observation y for each data point $\mathbf{x} = (x_0, x_1, \dots, x_{d-1}) \in \mathbb{R}^d$. The **linear model** $\mathbf{x} \mapsto \hat{y}$ approximates $\mathbf{x} \mapsto y$.¹

¹Should consider the output y here to be a random variable, with distribution that depends on \mathbf{x} . In simple linear regression, ε is a random variable and $y = p_0x_0 + p_1 + \varepsilon$.

Close only counts in horseshoes ...Regression

In linear regression, on indpt. variables x_0, x_1, \dots, x_{d-1} , had (affine) linear function $\hat{y} = p_0x_0 + p_1x_1 + \dots + p_{d-1}x_{d-1} + p_d$; values of function \leftrightarrow prediction \hat{y} ; error term ε , so that $y = \hat{y} + \varepsilon$. An observation y for each data point $\mathbf{x} = (x_0, x_1, \dots, x_{d-1}) \in \mathbb{R}^d$. The **linear model** $\mathbf{x} \mapsto \hat{y}$ approximates $\mathbf{x} \mapsto y$.¹ We expect that $|y - \hat{y}|$ is almost never exactly 0; a good model: one where $|y - \hat{y}|$ is small on average (but, still positive).

“Regression”

¹Should consider the output y here to be a random variable, with distribution that depends on \mathbf{x} . In simple linear regression, ε is a random variable and $y = p_0x_0 + p_1 + \varepsilon$.

Close only counts in horseshoes ...Regression

In linear regression, on indpt. variables x_0, x_1, \dots, x_{d-1} , had (affine) linear function $\hat{y} = p_0x_0 + p_1x_1 + \dots + p_{d-1}x_{d-1} + p_d$; values of function \leftrightarrow prediction \hat{y} ; error term ε , so that $y = \hat{y} + \varepsilon$. An observation y for each data point $\mathbf{x} = (x_0, x_1, \dots, x_{d-1}) \in \mathbb{R}^d$. The **linear model** $\mathbf{x} \mapsto \hat{y}$ approximates $\mathbf{x} \mapsto y$.¹ We expect that $|y - \hat{y}|$ is almost never exactly 0; a good model: one where $|y - \hat{y}|$ is small on average (but, still positive).

“Regression”

“Classification” tasks: the value y is a label and might not even be a number. The prediction \hat{y} is simply wrong, or not; close doesn't count. Good model: when $\hat{y} = y$ as often as possible.

¹Should consider the output y here to be a random variable, with distribution that depends on \mathbf{x} . In simple linear regression, ε is a random variable and $y = p_0x_0 + p_1 + \varepsilon$.

Outline

Classification tasks

Half-space model

Perceptron algorithm

A linear model for classification

Assume data is from \mathbb{R}^d for some $d > 0$ and we only have two labels (e.g., this is Spam (S) or it is Not spam (N)).

A linear model for classification

Assume data is from \mathbb{R}^d for some $d > 0$ and we only have two labels (e.g., this is Spam (S) or it is Not spam (N)).

A hyperplane in \mathbb{R}^d is an (affine) linear subspace that separates \mathbb{R}^d in two. Perhaps we get lucky and can find a hyperplane H so that data points with label S are on one side of H and data with label N are on the other side.

A linear model for classification

Assume data is from \mathbb{R}^d for some $d > 0$ and we only have two labels (e.g., this is Spam (S) or it is Not spam (N)).

A hyperplane in \mathbb{R}^d is an (affine) linear subspace that separates \mathbb{R}^d in two. Perhaps we get lucky and can find a hyperplane H so that data points with label S are on one side of H and data with label N are on the other side. Using coordinates (x_1, x_2, \dots, x_d) in \mathbb{R}^d , a hyperplane H may be determined from $d + 1$ numbers w_1, w_2, \dots, w_d , and b . It consists of solutions to

$$w_1x_1 + w_2x_2 + \dots + w_dx_d + b = 0.$$

- ▶ Rewriting in vector form: $\mathbf{w} = (w_1, w_2, \dots, w_d)$, look for solutions $\mathbf{x} \in \mathbb{R}^d$ to the equation $\mathbf{w} \cdot \mathbf{x} + b = 0$.
- ▶ \mathbf{w} is a vector that is orthogonal to a $(d - 1)$ -dimensional subspace of \mathbb{R}^d ; $|b|$ corresponds to a translation away from the origin.

Half-space model, continued

Using the notation from last slide:

a half-space model in \mathbb{R}^d is determined by $d + 1$ parameters

w_1, w_2, \dots, w_d, b ; the first d parameters grouped into a vector:

$\mathbf{w} = (w_1, w_2, \dots, w_d)$.

Half-space model, continued

Using the notation from last slide:

a half-space model in \mathbb{R}^d is determined by $d + 1$ parameters w_1, w_2, \dots, w_d, b ; the first d parameters grouped into a vector: $\mathbf{w} = (w_1, w_2, \dots, w_d)$.

Given $\mathbf{x} \in \mathbb{R}^d$, the side of the hyperplane it is on is determined by the sign of $\mathbf{w} \cdot \mathbf{x} + b$.

- ▶ (Positive side) Say that $h(\mathbf{x}) = 1$ if $\mathbf{w} \cdot \mathbf{x} + b > 0$.
- ▶ (Negative side) Say that $h(\mathbf{x}) = -1$ if $\mathbf{x} \cdot \mathbf{x} + b < 0$.

Half-space model, continued

Using the notation from last slide:

a half-space model in \mathbb{R}^d is determined by $d + 1$ parameters w_1, w_2, \dots, w_d, b ; the first d parameters grouped into a vector: $\mathbf{w} = (w_1, w_2, \dots, w_d)$.

Given $\mathbf{x} \in \mathbb{R}^d$, the side of the hyperplane it is on is determined by the sign of $\mathbf{w} \cdot \mathbf{x} + b$.

- ▶ (Positive side) Say that $h(\mathbf{x}) = 1$ if $\mathbf{w} \cdot \mathbf{x} + b > 0$.
- ▶ (Negative side) Say that $h(\mathbf{x}) = -1$ if $\mathbf{x} \cdot \mathbf{x} + b < 0$.

If there exists a hyperplane, given by some \mathbf{w}, b , so that \mathbf{x} has one of the labels if and only if it is on the positive side, the labeled data are called **linearly separable**.

Linearly separable




Figure: The hyperplane $H = \{(x, y) \in \mathbb{R}^2 : x + y + 1 = 0\}$, corresponding positive and negative regions, $\mathbf{w} = (1, 1)$, $b = 1$

Not linearly separable




Figure: A data set in \mathbb{R}^2 that is not linearly separable.

- A criterion (checkable, in theory) that is equivalent to “not linearly separable”?

Outline

Classification tasks

Half-space model

Perceptron algorithm

Setup for Perceptron algorithm

Assuming that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ is linearly separable, the Perceptron algorithm is a procedure that is guaranteed to find a hyperplane that separates the data.

Setup for Perceptron algorithm

Assuming that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ is linearly separable, the Perceptron algorithm is a procedure that is guaranteed to find a hyperplane that separates the data.

To describe it: for each \mathbf{x}_i , use X_i to denote the $(d + 1)$ -vector consisting of \mathbf{x}_i with 1 appended at the end;

Additionally, use W to denote the vector \mathbf{w} with b appended at the end.

Setup for Perceptron algorithm

Assuming that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ is linearly separable, the Perceptron algorithm is a procedure that is guaranteed to find a hyperplane that separates the data.

To describe it: for each \mathbf{x}_i , use X_i to denote the $(d + 1)$ -vector consisting of \mathbf{x}_i with 1 appended at the end;

Additionally, use W to denote the vector \mathbf{w} with b appended at the end.

Note that $W \cdot X_i = \mathbf{w} \cdot \mathbf{x}_i + b$.

For linearly separable data, our goal is to find $W \in \mathbb{R}^{d+1}$ so that, for all $1 \leq i \leq n$, $W \cdot X_i$ and y_i have the same sign (are both positive or both negative).

- Equivalently, we need $y_i W \cdot X_i > 0$ for all $1 \leq i \leq n$.

Perceptron algorithm


Suppose the data is linearly separable. Also, x is an $n \times d$ array of points, with i^{th} row equal to x_i , and y is array of the labels. The Perceptron algorithm finds W iteratively as follows.

```
input:  $x, y$   ##  $x$  is  $n$  by  $d$ ,  $y$  is  $1d$  array
 $X \leftarrow$  append 1 to each row of  $x$ 
 $W \leftarrow (0, 0, \dots, 0)$   ## Initial  $W$ 
while (exists  $i$  with  $y[i] * \text{dot}(W, X[i]) \leq 0$ ) {
     $W \leftarrow W + y[i] * X[i]$ 
}
return  $W$ 
```

Example

A simple example in \mathbb{R}^2 , with $n = 4$ points.

$$\mathbf{x}: \begin{bmatrix} -1 & 3 \\ -1 & -1 \\ 3 & -1 \\ 0 & 1.5 \end{bmatrix} \quad \mathbf{y}: \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$



Example, continued

A simple example in \mathbb{R}^2 , with $n = 4$ points.

$$\mathbf{x}: \begin{bmatrix} -1 & 3 \\ -1 & -1 \\ 3 & -1 \\ 0 & 1.5 \end{bmatrix} \qquad \mathbf{y}: \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

Beginning with $W^{(1)} = (0, 0, 0)$, next step:

$$W^{(2)} = \vec{0} + y_1 X_1 = -1 * (-1, 3, 1) = (1, -3, -1).$$

Next: since $y_1 W^{(2)} \cdot X_1 > 0$, check

$$y_2 W^{(2)} \cdot X_2 = -1 * (-1 + 3 - 1) = -1. \text{ So,}$$

$$W^{(3)} = W^{(2)} + y_2 X_2 = (2, -2, -2).$$

Continue in this way – on each step check dot products (in order) with

$y_1 X_1, y_2 X_2, y_3 X_3, y_4 X_4$. Eventually you return the vector

$$W^{(10)} = (4, -0.5, 1).$$

i.e., $H = \{(x_1, x_2) \in \mathbb{R}^2 : 4x_1 - 0.5x_2 + 1 = 0\}$ separates the points.

Perceptron algorithm, stopping time

Under our assumptions for Perceptron algorithm, a guarantee on eventually stopping.

Theorem

Define $R = \max_i |X_i|$ and $B = \min_i \{|V| : \forall i, y_i V \cdot X_i \geq 1\}$. Then, the Perceptron algorithm stops after at most $(RB)^2$ iterations and, when it stops with output W , then $y_i W \cdot X_i > 0$ for all $1 \leq i \leq n$.

Perceptron algorithm, stopping time

Under our assumptions for Perceptron algorithm, a guarantee on eventually stopping.

Theorem

Define $R = \max_i |X_i|$ and $B = \min_i \{|V| : \forall i, y_i V \cdot X_i \geq 1\}$. Then, the Perceptron algorithm stops after at most $(RB)^2$ iterations and, when it stops with output W , then $y_i W \cdot X_i > 0$ for all $1 \leq i \leq n$.

Idea of proof: Write W^* for vector that realizes the minimum B . Also, write $W^{(t)}$ for the vector W on the t^{th} step, with $W^{(1)} = (0, 0, \dots, 0)$.

Perceptron algorithm, stopping time

Under our assumptions for Perceptron algorithm, a guarantee on eventually stopping.

Theorem

Define $R = \max_i |X_i|$ and $B = \min_i \{|V| : \forall i, y_i V \cdot X_i \geq 1\}$. Then, the Perceptron algorithm stops after at most $(RB)^2$ iterations and, when it stops with output W , then $y_i W \cdot X_i > 0$ for all $1 \leq i \leq n$.

Idea of proof: Write W^* for vector that realizes the minimum B . Also, write $W^{(t)}$ for the vector W on the t^{th} step, with $W^{(1)} = (0, 0, \dots, 0)$. Using how $W^{(t+1)}$ is obtained from $W^{(t)}$, can show that $W^* \cdot W^{(T+1)} \geq T$ after $T + 1$ iterations. Also, using the condition on $W^{(T)}$ that necessitates an update, can show that $|W^{(T+1)}| \leq \sqrt{T}R$. (Both statements, use induction.)

Perceptron algorithm, stopping time

Under our assumptions for Perceptron algorithm, a guarantee on eventually stopping.

Theorem

Define $R = \max_i |X_i|$ and $B = \min_i \{|V| : \forall i, y_i V \cdot X_i \geq 1\}$. Then, the Perceptron algorithm stops after at most $(RB)^2$ iterations and, when it stops with output W , then $y_i W \cdot X_i > 0$ for all $1 \leq i \leq n$.

Idea of proof: Write W^* for vector that realizes the minimum B . Also, write $W^{(t)}$ for the vector W on the t^{th} step, with $W^{(1)} = (0, 0, \dots, 0)$. Using how $W^{(t+1)}$ is obtained from $W^{(t)}$, can show that $W^* \cdot W^{(T+1)} \geq T$ after $T + 1$ iterations. Also, using the condition on $W^{(T)}$ that necessitates an update, can show that $|W^{(T+1)}| \leq \sqrt{T}R$. (Both statements, use induction.)

Now, by Cauchy-Schwarz inequality, $T \leq BR\sqrt{T}$, which we can rearrange to $T \leq (BR)^2$.

Another example, the Iris data set

First discussed by R.A. Fisher in a 1936 paper, Iris data set commonly used in explanations. It contains 150 points in \mathbb{R}^4 , each for an individual iris flower from one of 3 species: *Iris setosa*, *Iris virginica*, and *Iris versicolor*. The 4 coordinates are measurements of sepal length, sepal width, petal length, and petal width (in cm).

Iris setosa points are linearly separable from the other two.

Labels: *Iris setosa* \leftarrow 1; *Other species* \leftarrow -1.

Begin by opening the notebook

`'perceptron-iris-notebook.ipynb'` ...After completing the algorithm, should get $\mathbf{w} = (1.3, 4.1, -5.2, -2.2)$ and $b = 1$.



Figure: Images by G. Robertson, E. Hunt, Radomil ©CC BY-SA 3.0