

A survey of some Machine Learning models

Chris Cornwell

April 1, 2025

Outline

Support Vector Machines

Neural Networks

Setup

Similar to a logistic model, a **support vector machine** is a model for binary classification, using a hyperplane, of the form $\{\mathbf{x} \in \mathbb{R}^d \mid \mathbf{w} \cdot \mathbf{x} + b = 0\}$, as decision boundary.

However, the optimization goal is different.

¹The multiple $\frac{1}{2}$ is non-essential here, but the convention is to include it.

²Recall, we pointed out how this is possible when the data is linearly separable.

Setup

Similar to a logistic model, a **support vector machine** is a model for binary classification, using a hyperplane, of the form $\{\mathbf{x} \in \mathbb{R}^d \mid \mathbf{w} \cdot \mathbf{x} + b = 0\}$, as decision boundary.

However, the optimization goal is different.

Given sample data $\mathcal{S} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, the goal is to minimize the value of $\frac{1}{2}|\mathbf{w}|^2$, the vector norm¹, subject to the condition that for all $1 \leq i \leq n$, $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ is satisfied.²

To work with a data set that is not linearly separable, one introduces so-called “slack variables” $\xi_i \geq 0, i = 1, \dots, n$ into the inequalities. They change to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1 - \xi_i$.

The reason for wanting to minimize $\frac{1}{2}|\mathbf{w}|^2$?

- ▶ Supposing no \mathbf{x}_i passes through hyperplane with parameters \mathbf{w}, b , we can scale both the normal vector and b so that $\min_i |\mathbf{w} \cdot \mathbf{x}_i + b| = 1$.
- ▶ The distance from any $\mathbf{x} \in \mathbb{R}^d$ to the hyperplane is $\frac{|\mathbf{w} \cdot \mathbf{x} + b|}{|\mathbf{w}|}$. So, if $\mathbf{x}_i \in \mathcal{S}$ is such that $|\mathbf{w} \cdot \mathbf{x}_i + b| = 1$, then its distance to decision boundary is $\rho = \frac{1}{|\mathbf{w}|}$.
- ▶ Want to maximize distance to decision boundary, so want to minimize $|\mathbf{w}|$.

¹The multiple $\frac{1}{2}$ is non-essential here, but the convention is to include it.

²Recall, we pointed out how this is possible when the data is linearly separable.

Constrained minimization and SVM

We can understand minimizing $\frac{1}{2}|\mathbf{w}|^2$ subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ with a Lagrangian. (Method of Lagrange multipliers; see Section 7.2 in the Mathematics for Machine Learning book.)

Constrained minimization and SVM

We can understand minimizing $\frac{1}{2}|\mathbf{w}|^2$ subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ with a Lagrangian. (Method of Lagrange multipliers; see Section 7.2 in the Mathematics for Machine Learning book.)

For $\underline{\alpha} = (\alpha_1, \dots, \alpha_n)$, with $\alpha_i \in \mathbb{R}$, Lagrangian is

$$L(\mathbf{w}, b, \underline{\alpha}) = \frac{1}{2}|\mathbf{w}|^2 - \sum_{i=1}^n \alpha_i (y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1) .$$

Constrained minimization and SVM

We can understand minimizing $\frac{1}{2}|\mathbf{w}|^2$ subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ with a Lagrangian. (Method of Lagrange multipliers; see Section 7.2 in the Mathematics for Machine Learning book.)

For $\underline{\alpha} = (\alpha_1, \dots, \alpha_n)$, with $\alpha_i \in \mathbb{R}$, Lagrangian is

$$L(\mathbf{w}, b, \underline{\alpha}) = \frac{1}{2}|\mathbf{w}|^2 - \sum_{i=1}^n \alpha_i (y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1).$$

It is minimized when

$$\nabla_{\mathbf{w}} L = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i;$$

$$\nabla_b L = 0 \quad \Rightarrow \quad \sum_{i=1}^n \alpha_i y_i = 0;$$

$$\alpha_i (y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1) = 0 \quad \Rightarrow \quad \alpha_i = 0 \quad \text{OR} \quad y_i(\mathbf{w} \cdot \mathbf{x}_i + b) = 1.$$

Constrained minimization and SVM

We can understand minimizing $\frac{1}{2}|\mathbf{w}|^2$ subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ with a Lagrangian. (Method of Lagrange multipliers; see Section 7.2 in the Mathematics for Machine Learning book.)

For $\underline{\alpha} = (\alpha_1, \dots, \alpha_n)$, with $\alpha_i \in \mathbb{R}$, Lagrangian is

$$L(\mathbf{w}, b, \underline{\alpha}) = \frac{1}{2}|\mathbf{w}|^2 - \sum_{i=1}^n \alpha_i (y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1).$$

It is minimized when

$$\nabla_{\mathbf{w}} L = 0 \quad \Rightarrow \quad \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i;$$

$$\nabla_b L = 0 \quad \Rightarrow \quad \sum_{i=1}^n \alpha_i y_i = 0;$$

$$\alpha_i (y_i(\mathbf{w} \cdot \mathbf{x}_i + b) - 1) = 0 \quad \Rightarrow \quad \alpha_i = 0 \quad \text{OR} \quad y_i(\mathbf{w} \cdot \mathbf{x}_i + b) = 1.$$

Support vectors are those \mathbf{x}_i for which $\alpha_i \neq 0$, and so $\mathbf{w} \cdot \mathbf{x}_i + b = \pm 1$.

Lagrangian Duality

Something interesting happens when we convert the previous Lagrangian optimization problem into its “Lagrangian dual problem.” This means that we take the minimum solution for \mathbf{w} , put it into $L(\mathbf{w}, b, \underline{\alpha})$ and want multipliers $\alpha_i \geq 0$ that *maximize* the value of this. That is, maximize

$$\frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right|^2 - \sum_{i=1}^n \alpha_i \left(y_i \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j \right) \cdot \mathbf{x}_i + y_i b - 1 \right).$$

Lagrangian Duality

Something interesting happens when we convert the previous Lagrangian optimization problem into its “Lagrangian dual problem.” This means that we take the minimum solution for \mathbf{w} , put it into $L(\mathbf{w}, b, \underline{\alpha})$ and want multipliers $\alpha_i \geq 0$ that *maximize* the value of this. That is, maximize

$$\frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right|^2 - \sum_{i=1}^n \alpha_i \left(y_i \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j \right) \cdot \mathbf{x}_i + y_i b - 1 \right).$$

Rearranged, you can rewrite it:

$$\max_{\underline{\alpha}} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j$$

subject to $\alpha_i \geq 0$ and $\sum_{i=1}^n \alpha_i y_i = 0$.

Lagrangian Duality

Something interesting happens when we convert the previous Lagrangian optimization problem into its “Lagrangian dual problem.” This means that we take the minimum solution for \mathbf{w} , put it into $L(\mathbf{w}, b, \underline{\alpha})$ and want multipliers $\alpha_i \geq 0$ that *maximize* the value of this. That is, maximize

$$\frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right|^2 - \sum_{i=1}^n \alpha_i \left(y_i \left(\sum_{j=1}^n \alpha_j y_j \mathbf{x}_j \right) \cdot \mathbf{x}_i + y_i b - 1 \right).$$

Rearranged, you can rewrite it:

$$\max_{\underline{\alpha}} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j$$

subject to $\alpha_i \geq 0$ and $\sum_{i=1}^n \alpha_i y_i = 0$.

This optimization problem only depends on knowing $\mathbf{x}_i \cdot \mathbf{x}_j$ for each (i, j) , and this leads to what are called **kernel methods** that are very computationally efficient and allow one to use SVM models that have non-linear decision boundaries.

SVMs via Gradient Descent

An alternative for optimizing an SVM classifier is to do so with a loss function. The loss function has a fair amount of similarity to the log-loss function we used in logistic regression; however, the per-example losses use a piecewise linear function.

SVMs via Gradient Descent

An alternative for optimizing an SVM classifier is to do so with a loss function. The loss function has a fair amount of similarity to the log-loss function we used in logistic regression; however, the per-example losses use a piecewise linear function.

When $y_i = 1$ then, writing $z_i = \mathbf{w} \cdot \mathbf{x}_i + b$, the per-example loss is $C \max(1 - z_i, 0)$ for some constant C . Call this $\text{Ccost}_1(z_i)$. When $y_i = -1$ (and so $\tilde{y}_i = 0$) then the per-example loss is $\text{Ccost}_0(z_i) = C \max(1 + z_i, 0)$.

SVMs via Gradient Descent

An alternative for optimizing an SVM classifier is to do so with a loss function. The loss function has a fair amount of similarity to the log-loss function we used in logistic regression; however, the per-example losses use a piecewise linear function.

When $y_i = 1$ then, writing $z_i = \mathbf{w} \cdot \mathbf{x}_i + b$, the per-example loss is $C \max(1 - z_i, 0)$ for some constant C . Call this $C \text{cost}_1(z_i)$. When $y_i = -1$ (and so $\tilde{y}_i = 0$) then the per-example loss is $C \text{cost}_0(z_i) = C \max(1 + z_i, 0)$. However, we also include the norm of \mathbf{w} in the loss:

$$\mathcal{L}_S(\mathbf{w}, b) = \frac{1}{2}|\mathbf{w}|^2 + \frac{1}{n} \sum_{i=1}^n C (\tilde{y}_i \text{cost}_1(\mathbf{w} \cdot \mathbf{x}_i + b) + (1 - \tilde{y}_i) \text{cost}_0(\mathbf{w} \cdot \mathbf{x}_i + b)) .$$

Outline

Support Vector Machines

Neural Networks

Single Layer

Fix a dimension d and a function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ (feel free to think of this as the logistic function, for now). For any $\omega = (\mathbf{w}, b) \in \mathbb{R}^{d+1}$, set

$$f_{\omega}(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b).$$

Single Layer

Fix a dimension d and a function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ (feel free to think of this as the logistic function, for now). For any $\omega = (\mathbf{w}, b) \in \mathbb{R}^{d+1}$, set

$$f_{\omega}(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b).$$

Given some collection $\omega_1 = (\mathbf{w}_1, b_1), \omega_2 = (\mathbf{w}_2, b_2), \dots, \omega_m = (\mathbf{w}_m, b_m)$, define $F : \mathbb{R}^d \rightarrow \mathbb{R}^m$ by setting

$$F(\mathbf{x}) = (f_{\omega_1}(\mathbf{x}), f_{\omega_2}(\mathbf{x}), \dots, f_{\omega_m}(\mathbf{x})).$$

Such a function F represents a single “layer map” of a neural network, with **activation function** σ .

Single Layer

Fix a dimension d and a function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ (feel free to think of this as the logistic function, for now). For any $\omega = (\mathbf{w}, b) \in \mathbb{R}^{d+1}$, set

$$f_{\omega}(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b).$$

Given some collection $\omega_1 = (\mathbf{w}_1, b_1), \omega_2 = (\mathbf{w}_2, b_2), \dots, \omega_m = (\mathbf{w}_m, b_m)$, define $F : \mathbb{R}^d \rightarrow \mathbb{R}^m$ by setting

$$F(\mathbf{x}) = (f_{\omega_1}(\mathbf{x}), f_{\omega_2}(\mathbf{x}), \dots, f_{\omega_m}(\mathbf{x})).$$

Such a function F represents a single “layer map” of a neural network, with **activation function** σ .

If we use W to denote the $m \times d$ matrix with \mathbf{w}_i as row i , and $\mathbf{b} = (b_1, \dots, b_m)$, then we can write $F(\mathbf{x}) = \sigma(W\mathbf{x} + \mathbf{b})$ (if we allow σ to be applied to a vector by applying the function to each component of the vector).

Single Layer

Fix a dimension d and a function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ (feel free to think of this as the logistic function, for now). For any $\omega = (\mathbf{w}, b) \in \mathbb{R}^{d+1}$, set

$$f_{\omega}(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b).$$

Given some collection $\omega_1 = (\mathbf{w}_1, b_1), \omega_2 = (\mathbf{w}_2, b_2), \dots, \omega_m = (\mathbf{w}_m, b_m)$, define $F : \mathbb{R}^d \rightarrow \mathbb{R}^m$ by setting

$$F(\mathbf{x}) = (f_{\omega_1}(\mathbf{x}), f_{\omega_2}(\mathbf{x}), \dots, f_{\omega_m}(\mathbf{x})).$$

Such a function F represents a single “layer map” of a neural network, with **activation function** σ .

If we use W to denote the $m \times d$ matrix with \mathbf{w}_i as row i , and $\mathbf{b} = (b_1, \dots, b_m)$, then we can write $F(\mathbf{x}) = \sigma(W\mathbf{x} + \mathbf{b})$ (if we allow σ to be applied to a vector by applying the function to each component of the vector).

W : called the “weight matrix” for the layer map; \mathbf{b} is called the “bias vector.”

Composing layers

A neural network is the result of composing some number of layer maps. That is, let m_0, m_1, \dots, m_L be positive integers and say that F_1, F_2, \dots, F_L are each layer maps (as in the previous slide), with $F_i : \mathbb{R}^{m_{i-1}} \rightarrow \mathbb{R}^{m_i}$.

- so, for each $1 \leq i \leq L$, there is a weight matrix W_i (which is $m_i \times m_{i-1}$) and bias vector $\mathbf{b}_i \in \mathbb{R}^{m_i}$.

Composing layers

A neural network is the result of composing some number of layer maps. That is, let m_0, m_1, \dots, m_L be positive integers and say that F_1, F_2, \dots, F_L are each layer maps (as in the previous slide), with $F_i : \mathbb{R}^{m_{i-1}} \rightarrow \mathbb{R}^{m_i}$.

- so, for each $1 \leq i \leq L$, there is a weight matrix W_i (which is $m_i \times m_{i-1}$) and bias vector $\mathbf{b}_i \in \mathbb{R}^{m_i}$.

The *fully connected* neural network, with activation function σ , associated to this collection of weight matrices and bias vectors is the parameterized function

$$f(\mathbf{x}) = F_L \circ F_{L-1} \circ \dots \circ F_2 \circ F_1(\mathbf{x}).$$

Composing layers


A neural network is the result of composing some number of layer maps. That is, let m_0, m_1, \dots, m_L be positive integers and say that F_1, F_2, \dots, F_L are each layer maps (as in the previous slide), with $F_i : \mathbb{R}^{m_{i-1}} \rightarrow \mathbb{R}^{m_i}$.

- so, for each $1 \leq i \leq L$, there is a weight matrix W_i (which is $m_i \times m_{i-1}$) and bias vector $\mathbf{b}_i \in \mathbb{R}^{m_i}$.

The *fully connected* neural network, with activation function σ , associated to this collection of weight matrices and bias vectors is the parameterized function

$$f(\mathbf{x}) = F_L \circ F_{L-1} \circ \dots \circ F_2 \circ F_1(\mathbf{x}).$$

The parameters for the function are $(W_1, \mathbf{b}_1, W_2, \mathbf{b}_2, \dots, W_L, \mathbf{b}_L)$ (flattened out as a vector). Layers 1, 2, \dots , $L - 1$ are often called “hidden layers,” m_i is the width of layer i , and L is the “depth” of the network.



Variety of tasks for Neural Networks

Depending on choices that are made, neural networks have the ability to be used for a very wide variety of tasks.

Variety of tasks for Neural Networks

Depending on choices that are made, neural networks have the ability to be used for a very wide variety of tasks.

- Regression tasks: Say that $m_L = 1$. By not using the activation function on the final layer map F_L (or using a different function), can find parameters to fit data for a regression task on points in \mathbb{R}^{m_0+1} .

Variety of tasks for Neural Networks

Depending on choices that are made, neural networks have the ability to be used for a very wide variety of tasks.

- ▶ Regression tasks: Say that $m_L = 1$. By not using the activation function on the final layer map F_L (or using a different function), can find parameters to fit data for a regression task on points in \mathbb{R}^{m_0+1} .
- ▶ Classification tasks, with any number of possible labels: In place of applying the activation function to last layer, use something called the **softmax** function - makes the output of $f(\mathbf{x})$ be a probability vector $\rightsquigarrow m_L$ different possible classes.

Variety of tasks for Neural Networks

Depending on choices that are made, neural networks have the ability to be used for a very wide variety of tasks.

- ▶ Regression tasks: Say that $m_L = 1$. By not using the activation function on the final layer map F_L (or using a different function), can find parameters to fit data for a regression task on points in \mathbb{R}^{m_0+1} .
- ▶ Classification tasks, with any number of possible labels: In place of applying the activation function to last layer, use something called the **softmax** function - makes the output of $f(\mathbf{x})$ be a probability vector $\rightsquigarrow m_L$ different possible classes.
- ▶ Finding a smaller set of “useful” coordinates for data: e.g. assigning to each word (or “token”) in a dictionary a vector in \mathbb{R}^{100} so that semantically similar words are assigned similar vectors. (Has $m_0 = m_L$ and m_i much smaller (100) for some hidden layer i ; this is a variation on something called an *Autoencoder*.)

Learning parameters of a neural network

Done with some method that is based on gradient descent, usually mini-batch gradient descent. When doing mini-batch gradient descent, all points in the training data are randomly grouped into a mini-batch S' (so the union of all of them is all the training data). Gradient descent is then done with each mini-batch; once every one has been used \rightarrow one “epoch” of training. You usually do many epochs of training, newly re-organizing the mini-batches with each epoch.

Learning parameters of a neural network

Done with some method that is based on gradient descent, usually mini-batch gradient descent. When doing mini-batch gradient descent, all points in the training data are randomly grouped into a mini-batch S' (so the union of all of them is all the training data). Gradient descent is then done with each mini-batch; once every one has been used \rightarrow one “epoch” of training. You usually do many epochs of training, newly re-organizing the mini-batches with each epoch.

The neural network function is built by composing functions. Computation of partial derivatives is efficiently handled by using the Chain rule – a certain implementation of the chain rule called *backpropagation*.

Learning parameters of a neural network

Done with some method that is based on gradient descent, usually mini-batch gradient descent. When doing mini-batch gradient descent, all points in the training data are randomly grouped into a mini-batch \mathcal{S}' (so the union of all of them is all the training data). Gradient descent is then done with each mini-batch; once every one has been used \rightarrow one “epoch” of training. You usually do many epochs of training, newly re-organizing the mini-batches with each epoch.

The neural network function is built by composing functions. Computation of partial derivatives is efficiently handled by using the Chain rule – a certain implementation of the chain rule called *backpropagation*.

The update step often involves something more than just learning rate times the gradient: a “momentum” term, and some 2nd derivative information.

