

Selecting the Machine Learning Model

Chris Cornwell

April 14, 2025

Outline

Overfitting

Having only concern be to minimize empirical loss

Outline

Overfitting

Having only concern be to minimize empirical loss

Intro

Given sample data \mathcal{S} , consisting of greyscale images of handwritten digits 0, 1, . . . , 9, say we want a decision tree model to predict the correct digit, given an image.

Let the images be p pixels by p pixels. For an image, we might take the p^2 greyscale values (0–255) and lay them end to end to get a vector in \mathbb{R}^{p^2} . Then, could train a decision tree with however many splits are needed so that each leaf has points in \mathcal{S} with only one label.

- ▶ Gives 100% accuracy on \mathcal{S} , the data used to determine the parameters.
- ▶ But, the model won't perform nearly as well on data not from \mathcal{S} .

Let's look at doing exactly this, using images of digits available from the Python package `scikit-learn`.

Decision tree model for handwritten digits

The images of digits from scikit-learn are 8 pixels by 8 pixels. With the submodule datasets imported from sklearn, the data can be loaded as follows.

```
1 | data = datasets.load_digits()  
2 | x = data.images  
3 | y = data.target
```

To convert each 8×8 array to a vector in \mathbb{R}^{64} , can use the reshape method; after assigning x and y as above, code below would output (1797, 64).

```
1 | # make each 8x8 array into 1d array with 64 entries  
2 | x = x.reshape(len(x), -1)  
3 | x.shape
```

Randomly select 20% of the data to test model with – it won't be used to determine the decision tree.

```
1 | n_test = int(0.2*len(x))  
2 | test_indices = np.random.choice(len(x), size=n_test, replace=False)  
3 | train_indices = np.delete(np.arange(len(x)), test_indices)  
4 | x_test, y_test = x[test_indices], y[test_indices]  
5 | x_train, y_train = x[train_indices], y[train_indices]
```

**Alternatively, you can use the function train_test_split from the scikit-learn submodule model_selection.*

Decision tree model for handwritten digits

With the setup from the last slide, use the training data to determine the decision tree.

There is a submodule `tree` within `scikit-learn` that we can use.

```
1 | model = tree.DecisionTreeClassifier()  
2 | model.fit(x_train, y_train)
```

The default behavior of the `.fit()` method is that the tree has as many splits (decision stumps) as needed so that each leaf consists of points with a single label. The accuracy on `x_train` is, therefore, 100%.

The accuracy of the model on `x_test` will depend some on which points were put into `x_test`; however, it tends to be around only 85%.

What happened is that the *hypothesis class*, that set of functions that were possible outcomes to be the trained model, was too large (allowed too many possibilities for the resulting model).

Keeping the hypothesis class small

We discuss a (somewhat absurd) example to demonstrate the point that
small empirical loss \nRightarrow small population loss

Allowing for *any* function as a possible model, for any training data $\mathcal{S} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, with labels in $\{1, -1\}$, we set

$$f_{\mathcal{S}}(\mathbf{x}) = \begin{cases} y_i, & \text{if } \mathbf{x} = \mathbf{x}_i \\ -1, & \text{otherwise.} \end{cases}$$

Regardless of \mathcal{S} , or the population distribution that it is a sample from, the empirical loss $\mathcal{L}_{\mathcal{S}}(f_{\mathcal{S}}) = \frac{\#\{i: y_i \neq f_{\mathcal{S}}(\mathbf{x}_i)\}}{n}$ is zero.




Figure: Upper-left square labeled -1

Keeping the hypothesis class small

We discuss a (somewhat absurd) example to demonstrate the point that
small empirical loss \nRightarrow small population loss

Suppose that the population data consists of points in $[0, 1]^2$ (the square of points with both coordinates between 0 and 1), and that a point (x_1, x_2) has label -1 if and only if $0 \leq x_1 < 0.5$ and $0.5 < x_2 \leq 1$ (See Figure).

Since a sample \mathcal{S} is finite, a randomly chosen point has probability 0 of being in \mathcal{S} . Thus, given a random $\mathbf{x} \in [0, 1]^2$, with probability $3/4$ the predicted label $f_{\mathcal{S}}(\mathbf{x})$ is wrong.

The issue with both this model, and the decision tree that had no restriction on depth, is that the **variance** of the hypothesis class is too large.

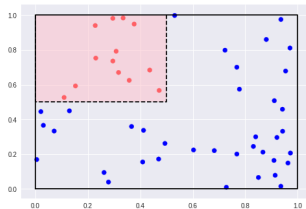


Figure: Upper-left square labeled -1

Bias-Variance Trade-off

Suppose that we have chosen a hypothesis class (a class of parameterized functions, say), and a procedure that, given a training sample \mathcal{S} from the population, determines a prediction function $f_{\mathcal{S}}$.

In order to understand the *expected* population loss, we want to consider the variance of $f_{\mathcal{S}}(\mathbf{x})$, not only over the input space but over possible training sets \mathcal{S} . That is, over all (\mathbf{x}, y) and \mathcal{S} , the expected value $\mathbb{E}[(f_{\mathcal{S}}(\mathbf{x}) - y)^2]$.

Let \bar{h} be the expected prediction function, over training sets.¹ In practice, you could approximate \bar{h} in the following way: for samples $\mathcal{S}_1, \dots, \mathcal{S}_N$, each with the same number of points, drawn i.i.d. from the population, then for sufficiently large N ,

$$\bar{h}(\mathbf{x}) \approx \frac{1}{N} \sum_{i=1}^N h_{\mathcal{S}_i}(\mathbf{x}).$$

¹This takes some advanced mathematics to carefully describe. It means there is a “measure” on sets of functions, which allows you to integrate over a set of functions to get a new function.

Bias-Variance Trade-off

$$\begin{aligned}\mathbb{E}_{S, \mathbf{x}, y}[(f_S(\mathbf{x}) - y)^2] \\&= \mathbb{E}_{S, \mathbf{x}, y}[(f_S(\mathbf{x}) - \bar{f}(\mathbf{x}) + \bar{f}(\mathbf{x}) - y)^2] \\&= \mathbb{E}_{S, \mathbf{x}}[(f_S(\mathbf{x}) - \bar{f}(\mathbf{x}))^2] + \mathbb{E}_{\mathbf{x}, y}[(\bar{f}(\mathbf{x}) - y)^2] + 2\mathbb{E}_{S, \mathbf{x}, y}[(f_S(\mathbf{x}) - \bar{f}(\mathbf{x}))(\bar{f}(\mathbf{x}) - y)]\end{aligned}$$

Since $\bar{f}(\mathbf{x}) = \mathbb{E}_S[f_S(\mathbf{x})]$, the last term vanishes and so

$$\mathbb{E}_{S, \mathbf{x}, y}[(f_S(\mathbf{x}) - y)^2] = \mathbb{E}_{S, \mathbf{x}}[(f_S(\mathbf{x}) - \bar{f}(\mathbf{x}))^2] + \mathbb{E}_{\mathbf{x}, y}[(\bar{f}(\mathbf{x}) - y)^2].$$

A similar argument shows that

$$\mathbb{E}_{\mathbf{x}, y}[(\bar{f}(\mathbf{x}) - y)^2] = \mathbb{E}_{\mathbf{x}}[(\bar{f}(\mathbf{x}) - \bar{y}(\mathbf{x}))^2] + \mathbb{E}_{\mathbf{x}, y}[(\bar{y}(\mathbf{x}) - y)^2].$$

And so

$$\mathbb{E}_{S, \mathbf{x}, y}[(f_S(\mathbf{x}) - y)^2] = \mathbb{E}_{S, \mathbf{x}}[(f_S(\mathbf{x}) - \bar{f}(\mathbf{x}))^2] + \mathbb{E}_{\mathbf{x}}[(\bar{f}(\mathbf{x}) - \bar{y}(\mathbf{x}))^2] + \mathbb{E}_{\mathbf{x}, y}[(\bar{y}(\mathbf{x}) - y)^2].$$

Variance \uparrow

Bias² \uparrow

Noise \uparrow